

● ENTERPRISE AI SECURITY GATEWAY

The AI Gateway built for *Zero Leaks*.

One integration. Any AI model. Complete data control, policy enforcement, and cost governance — for every employee, every prompt.

EU AI ACT

GDPR

SOC 2

ISO 27001

ISO 42001

NIST AI RMF

CUSTOMER DEMO • 2026

THE REALITY TODAY

Your people are already using AI. *Without you.*

Employees paste code, contracts, customer records, and financials into ChatGPT, Claude, and Gemini every day — through personal accounts, browser tabs, and unmanaged tools. You have no visibility, no policy, and no record.

1 in 5

files pasted into AI contain sensitive data

0

audit trail on shadow AI usage

∞

once data leaves, it can't be recalled

€35M

max EU AI Act penalty exposure

Six gaps every enterprise has right now.

LEAK

Sensitive data leaks to AI models

Credentials, PII, source code, and financials flow straight to third-party models — irreversibly.

CONTROL

No enterprise controls

No way to say which teams use which models for what. Every employee is an open door.

POLICY

No governance policy

No enforced rules, no versioning, no evidence for auditors or regulators.

LICENSES

Fragmented model licenses

Separate contracts for OpenAI, Anthropic, Google — duplicated spend, no leverage.

WASTE

Token waste, no cost control

Bloated prompts and redundant context burn budget with zero visibility per team or user.

ADOPTION

No effective, safe way to use AI

Employees want to use AI well; without guardrails, IT can only say "no."

WHAT IT COSTS YOU

"No" is not a strategy.
Neither is *"hope."*

Block AI entirely

Employees route around IT with personal accounts. You lose the productivity *and* the visibility. Shadow AI grows in the dark.

Allow AI freely

Every prompt is an uncontrolled, unaudited data-exfiltration path. One paste becomes a breach, a fine, or a headline.

There is a third option: **govern it.**

THE SOLUTION

One gateway between your people and every *AI* model.

Kotwal sits in the middle. Nothing sensitive leaves without passing policy. Every interaction is inspected, logged, and controlled — while your people keep working exactly as they do today.

INSPECT

REDACT

ENFORCE

ROUTE

AUDIT


OPTIMIZE

Every prompt inspected at the gate.



DATA LEAK PREVENTION

A real key never reaches the model.



SMART REDACTION

The AI still helps — but never sees *who*.

WHAT THE USER TYPED

Write a follow-up to `Jennifer Martinez` at
`jennifer.martinez@northwind-capital.com` about the Q1
renewal.

WHAT THE MODEL RECEIVES

Write a follow-up to `<PERSON_NAME>` at `<EMAIL>` about the Q1
renewal.

Identity is tokenized on the way out and restored in the reply your rep sees. The model does its job; the data stays inside your boundary.

Nine data classes.

One decision, *every prompt*.

BLOCK**Credentials**

API keys, tokens, private keys, DB connection strings.

BLOCK**Prompt injection**

Jailbreaks, system-override & indirect injection attempts.

BLOCK**Dangerous output & IP**

Destructive SQL/shell, malicious code, proprietary identifiers.

REDACT**Financial**

Cards (Luhn), IBAN, sort codes, account numbers.

REDACT**Government IDs**

SSN, Aadhaar, PAN, CPF, DNI — checksum-validated.

REDACT**Personal data**

Names, emails, phone numbers, physical addresses.

Plus network/infra references and configurable topic restrictions — **logged, warned, or blocked** to your policy.

POLICY, NOT GUESSWORK

One engine. *Your rules.*

EU

GDPR

Tighter thresholds for personal data across the EU.

HEALTH

HIPAA

PHI & personal identifiers escalated to hard block.

PAYMENTS

PCI-DSS

Financial data blocked outright, not just redacted.

INDIA


DPDP

India Digital Personal Data Protection preset.

Pick a jurisdiction preset or tune per team and role — every decision is versioned and audit-logged.

Everything, in one dashboard.

AI USAGE BY DEPARTMENT · THIS MONTH



Blocked prompts

128 THIS WEEK


Redacted prompts

342 THIS WEEK

Clean prompts

18,905 THIS WEEK

ENFORCEMENT OUTCOMES



clean · redacted · blocked

Policy version

v2 · GDPR

EVERYTHING INCLUDED

Controlled, safe, and ready to scale.

- ✓ Data-leak prevention on every prompt (credentials, PII, financials, IP)
- ✓ Multi-model gateway — one integration for OpenAI, Anthropic, Google & more
- ✓ Policy enforcement per team, role, and data category
- ✓ Full audit log of every interaction — evidence-ready
- ✓ Usage analytics & cost control across every provider
- ✓ Input-token optimization — 20–40% cost reduction
- ✓ Custom agents grounded on your own knowledge base
- ✓ Role-based access, project workspaces, VS Code plugin

Consolidate licenses. Cut the *token waste*.

CONSOLIDATE

One contract, all models

Replace scattered per-provider licenses with a single governed gateway and real negotiating leverage.

OPTIMIZE

20–40% fewer tokens

Automatic prompt compression and context trimming cut redundant tokens with no quality loss.

BUDGET

Per-team budgets

Prepaid pools, hard/soft caps, and low-balance alerts — spend visibility down to the user.

Governance that **pays for itself**.

YOUR DATA, YOUR INFRASTRUCTURE

Run it wherever your data must live.

SAAS

Our cloud

Zero ops. Encrypted at rest with your key.

VPC

Self-hosted

Your VPC, your DB, your keys. We never see the data.

MANAGED

Private cloud

Your infra, we operate it via secure access.

AIR-GAP

On-prem

Fully isolated for the highest-classification environments.

Same product, everywhere. In the self-hosted model there is nothing to trust — the data never leaves your walls.

COMPLIANCE BUILT IN

Evidence-ready for every framework.

100%

prompt coverage — no blind spots

<50ms

enforcement overhead

6

frameworks mapped

AES-256

GCM encryption at rest

EU AI ACT

GDPR

HIPAA

PCI-DSS

SOC 2

ISO 27001

ISO 42001

NIST AI RMF

DPDP (INDIA)

 GET STARTED

See Kotwal stop a real leak — *live*.

A 30-minute walkthrough on your own AI risk scenarios. See the gateway, policy enforcement, and the dashboard in action.

NO COMMITMENT

30 MINUTES

YOUR DATA SCENARIOS

AIKOTWAL.COM · ENTERPRISE AI SECURITY GATEWAY