

The control layer for *enterprise AI*.

Kotwal is the gateway every company will need between its employees and AI — stopping data leaks, enforcing policy, and controlling cost across every model, from one place.

INVESTOR OVERVIEW · CONFIDENTIAL

THE ONE LINE

Employees → Kotwal → any AI model

Every prompt inspected, redacted, and governed in under 50ms — invisibly. Safety that adoption doesn't fight.

PROTECT

Zero leaks

CONTROL

One dashboard

WHY NOW

AI adoption outran AI control.

Every enterprise now has thousands of employees using AI daily — and almost none has a way to govern it. The tools arrived overnight; the guardrails didn't. Regulation (EU AI Act) is landing at the same moment. This gap is being created faster than anyone can close it.

2026

EU AI Act enforcement begins — fines to €35M

78%

of employees use AI tools their employer can't see

Now

the window to become the default control layer

THE PROBLEM

Enterprises are flying blind on their biggest new risk.

Data walks out the door

Employees paste secrets, customer data, and IP into AI tools. Once it's sent, it's gone — no recall, no record.

No policy, no control

No enforced rules on who uses which model for what. IT's only lever today is a blanket "no" — which nobody obeys.

Cost sprawl

Separate licenses per provider, runaway token spend, zero visibility. AI budgets balloon with no accountability.

The choice today is **block AI and lose the productivity**, or **allow it and lose the data**. Both are wrong.

THE SOLUTION

One gateway. Every model.

Total control.

Kotwal sits between employees and every AI provider. It inspects each prompt, blocks or redacts sensitive data in real time, enforces policy per team, routes across models, and gives finance one dashboard for cost — without changing how anyone works.

PROTECT

Stop leaks

Real-time detection + redaction.

GOVERN

Enforce policy

Per team, role, and model.

ROUTE

Any model

One integration, all providers.

CONTROL

Cut cost

Budgets + token optimization.

Governance that adds zero friction.

Others force a trade-off: security *or* usability. Kotwal's engine inspects and redacts in under 50ms, invisibly — the employee gets their answer, the sensitive data never leaves. That's the unlock: safety people don't fight, so adoption actually happens.

SPEED

<50ms inline

Detection and redaction fast enough to sit on every prompt, invisibly.

BREADTH

One layer, everything

DLP, policy, multi-model routing, cost, agents, audit — unified, not bolted together.


TRUST

Runs in your VPC

Self-hostable — the data never leaves the customer's walls.
Removes the buyer's #1 objection.

IN ONE PICTURE

Every prompt passes the gate.



SENSITIVE DATA STOPPED OR MASKED BEFORE IT EVER REACHES A MODEL.

A category forming at *regulatory speed*.

TAM

\$60B+ · AI security & governance by 2030

SAM

\$18B · regulated & data-sensitive enterprises

SOM

\$1.5B · beachhead: finance, legal,
health

Every enterprise adopting AI needs this layer. Regulation turns it from optional to mandatory.

MARKET DRIVER

Enterprise AI spend

\$300B+ by 2030

Governance is the tax on all of it — and nobody owns that layer yet.

GROWTH

Category CAGR

~35% / yr

Fastest growing slice of enterprise security spend

WHY THE TIMING IS EXACT

The numbers behind the urgency.

\$300B+

projected enterprise AI spend by 2030 — every dollar needs a control layer

~35%

annual growth in AI security & governance — outpacing all other security categories

€35M

max EU AI Act penalty — governance becomes a board-mandated purchase in 2026

78%

of employees already use AI tools their employer cannot see or control

1 in 5

AI prompts in the enterprise contain sensitive or regulated data

20–40%

token cost reduction Kotwal delivers — the product pays for itself

Recurring, expanding, and sticky.

PER SEAT

Subscription

Priced per active user, per month. Grows automatically as the customer rolls AI out wider.

USAGE

Model spend

Gateway sits on all AI traffic — a natural, expanding usage layer on top of seats.

STICKY

Control plane


Once every prompt, policy, and audit trail runs through Kotwal, switching cost is enormous.

Land with security, expand with cost savings and adoption — the product pays for itself via token optimization alone.

WHY WE WIN

The default control layer compounds.

- ✓ **Data moat:** every enterprise's policies, audit trail, and detection tuning live in Kotwal.
- ✓ **Integration moat:** the gateway becomes plumbing across every AI tool and workflow.
- ✓ **Timing moat:** becoming the standard before the category consolidates.
- ✓ **Trust moat:** self-hosted deployment wins the most security-conscious buyers first.



ILLUSTRATIVE GROWTH — PIPELINE & DEPLOYMENTS

Founder Name

CO-FOUNDER & CEO

[Add 4–6 lines: background and prior companies, domain expertise in security / enterprise software, a notable achievement or exit, and the specific reason this founder is uniquely positioned to build Kotwal. There is room here for a full short bio — write it as you'd say it to an investor, without crowding the card.]

PHOTO — ADD LATER
(LANDSCAPE OR SQUARE)

Founder Name

CO-FOUNDER & CTO

[Add 4–6 lines: engineering background, deep systems / security / ML experience, notable prior work or scale shipped, and the technical edge behind Kotwal's sub-50ms detection and multi-model gateway. Enough space for a substantive bio that establishes real technical credibility.]

Raising to own the *control layer*.

We're raising to accelerate go-to-market, deepen the detection engine, and land the enterprises defining this category — before anyone else does.

PRODUCT

Engine & scale

Deepen context-aware detection and enterprise deployment.

GTM

Land & expand

Sales into regulated, data-sensitive enterprises.

TEAM

Key hires

Security, ML, and enterprise sales talent.